

# Downloading files with curl

*How to download files straight from the command-line interface*

The `curl` tool lets us fetch a given URL from the command-line. Sometimes we want to save a web file to our own computer. Other times we might pipe it directly into another program. Either way, `curl` has us covered.

See its [documentation here](#).

This is the basic usage of `curl` :

```
curl http://some.url --output some.file
```

That `--output` flag denotes the filename ( `some.file` ) of the downloaded URL ( `http://some.url` )

Let's try it with a basic website address:

```
curl http://example.com --output my.file
```

Besides the display of a progress indicator (which I explain below), you don't have much indication of what `curl` actually downloaded. So let's confirm that a file named `my.file` was actually downloaded.

Using the `ls` command will show the contents of the directory:

```
ls
```

Which outputs:

```
my.file
```

And if you use `cat` to output the contents of `my.file`, like so:

```
cat my.file
```

– you will the HTML that powers <http://example.com>

## I thought Unix was supposed to be quiet?

Let's back up a bit: when you first ran the `curl` command, you might have seen a quick blip of a progress indicator:

```

% Total      % Received % Xferd  Average Speed
Time        Time
                                Dload  Upload
Total      Spent
100 1270 100 1270    0      0 50573      0 -
-:--:--  -:--:--

```

If you remember the [Basics of the Unix Philosophy](#), one of the tenets is:

Rule of Silence: When a program has nothing surprising to say, it should say nothing.

In the example of `curl`, the author apparently believes that it's important to tell the user the progress of the download. For a very small file, that status display is not terribly helpful. Let's try it with a bigger file (this is the baby names file from the [Social Security Administration](#)) to see how the progress indicator animates:

```

curl http://stash.compciv.org/ssa_baby_names/names.zip \
  --output babynames.zip

```

**Quick note:** If you're new to the command-line, you're probably used to commands executing every time you hit **Enter**. In this case, the command is so long (because of the URL) that I broke it down into two lines with the use of the **backslash**, i.e. `\`

This is *solely to make it easier for you to read*. As far as the computer cares, it just joins the two lines together as if that backslash weren't there and runs it as one command.

## Make curl silent

The `curl` progress indicator is a nice affordance, but let's just see if we get `curl` to act like all of our Unix tools. In `curl`'s [documentation of options](#), there is an option for silence:

```
-s, --silent
```

Silent or quiet mode. Don't show progress meter or error messages. Makes Curl mute. It will still output the data you ask for, potentially even to the terminal/stdout unless you redirect it.

Try it out:

```
curl http://example.com --output my.file --silent
```

## Repeat and break things

So those are the basics for the `curl` command. There are many, many more options, but for now, we know how to use `curl` to do something that is actually quite powerful: fetch a file, anywhere on the Internet, from the simple confines of our command-line.

Before we go further, though, let's look at the various ways this simple command can be re-written and, more crucially, screwed up:

## Shortened options

As you might have noticed in the `--silent` documentation, it lists the alternative form of `-s`. Many options for many tools have a shortened alias. In fact, `--output` can be shortened to `-o`

```
curl http://example.com -o my.file -s
```

**Now watch out:** the number of hyphens is **not** something you can mess up on; the following commands would cause an error or other unexpected behavior:

```
curl http://example.com -o my.file -silent
```

```
curl http://example.com -output my.file -s
```

```
curl http://example.com --o my.file --s
```

Also, mind the position of `my.file`, which can be thought of as the **argument** to the `-o` **option**.

The **argument** *must* follow after the `-o` ...because `curl`.

If you instead executed this:

```
curl http://example.com -o -s my.file
```

How would `curl` know that `my.file`, and not `-s` is the **argument**, i.e. what you want to name the content of the downloaded URL?

In fact, you might see that you've created a file named `-s` ...which is not the end of the world, but not something you want to happen unwittingly.

## Order of options

By and large (from what I can think of at the top of my head), the *order* of the options doesn't matter:

```
curl http://example.com -s -o my.file
```

In fact, the URL, `http://example.com`, can be placed anywhere in the mix:

```
curl -s http://example.com -o my.file
```

```
curl -s -o my.file http://example.com
```

A couple of things to note:

1. The way that the URL, what you might consider the main **argument** for the `curl` command, can be placed anywhere after the command is *not* the way that *all* commands have been designed. So it always pays to read the documentation with every new command.
2. Notice how `-s http://example.com` doesn't cause a problem. That's because the `-s` option doesn't take an argument. But try the following:

```
curl -s -o http://example.com my.file
```

And you *will* have a problem.

## No options at all

The last thing to consider is what happens when you just `curl` for a URL with no options (which, after all, should be *optional*). Before you try it, think about another part of the Unix philosophy:

This is the Unix philosophy: Write programs that do one thing and do it well. Write programs to work together. **Write programs to handle text streams, because that is a universal interface.**

If you `curl` without any options except for the URL, the content of the URL (whether it's a webpage, or a binary file, such as an image or a zip file) will be printed out to screen. Try it:

```
curl http://example.com
```

Output:

```
<!doctype html>
<html>
<head>
  <title>Example Domain</title>

  <meta charset="utf-8" />
  <meta http-equiv="Content-type" content="text/html; charset=utf-8" />
  <meta name="viewport" content="width=device-width, initial-scale=1" />

... (and so forth)
```

Even with the small amount of HTML code that makes up the `http://example.com` webpage, it's too much for human eyes to process (and reading raw HTML wasn't meant for humans).



# Standard output and connecting programs

But what if we wanted to send the contents of a web file to *another* program? Maybe to `wc`, which is used to count words and lines? Then we can use the powerful Unix feature of **pipes**. In this example, I'm using `curl`'s silent option so that only the output of `wc` (and not the progress indicator) is seen. Also, I'm using the `-l` option for `wc` to just get the number of *lines* in the HTML for example.com:

```
curl -s http://example.com | wc -l
```

Number of lines in `example.com` is: `50`

Now, you could've also done the same in two lines:

```
curl -s http://example.com -o temp.file  
wc -l temp.file
```

But not only is that less *elegant*, it also requires *creating* a new file called `temp.file`. Now, this is a trivial concern, but someday, you may work with systems and data flows in which temporarily saving a file is not an available luxury (think of massive files).